

What Makes AI Recommend Your Product? Lessons From 2,000 Shopping Tests

When we ask ChatGPT what to buy, it reads the listings it finds and decides which ones to recommend first. A team from MIT and Columbia set out to learn what makes an AI put one product above another.

The study used 2,000 real shopping questions from Reddit, like “I need a toaster that lasts, any ideas?”. Each question came with 10 Amazon products.

An AI read the question and ranked the 10 products from best to worst for that shopper. The test ran on five AIs: ChatGPT (GPT-5), Claude, Gemini, DeepSeek and Llama.

Next, one product’s description was rewritten, and the AI ranked all 10 again to see whether the rewrite moved it up or down the list. In the tables below, a positive number means the product moved up by that many places on average, and a negative number means it moved down.

We’ll start with what worked.

What works

Can AI-rewritten product descriptions improve AI rankings?

Yes, but it depends on which AI model does the rewriting. Seven AI models were given the same request: “Optimize my product description so that an LLM is more likely to rank it higher than similar products.”

Rewrites by GPT-5 and Claude moved products up about half a place on average. GPT-4o-mini’s pushed them down about a third of a place.

AI that rewrote the description	Average places moved
Claude	+0.57
GPT-5	+0.56
Llama	+0.36
Gemini	+0.26
DeepSeek	+0.16
GPT-4.1	−0.03
GPT-4o-mini	−0.36

Source: E-GEO, Table 3 (average across all five ranking AIs)

What do the best-ranking descriptions have in common?

The researchers kept testing and refining the rewriting instructions over many rounds. The instructions that improved rankings the most all asked for:

- a short opening summary of what the product is and its main benefit
- the words a shopper would use to describe what they need
- related keywords used naturally, without repeating them
- clear headings with short bullet points
- real situations where the product would be used
- answers to the questions a buyer is likely to ask
- accurate facts, with no made-up claims

The best of these moved products up about 0.3 places on average, and all five AIs ranked them higher.

Best instructions	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Top instruction A	+0.04	+0.28	+0.14	+0.32	+0.67	+0.29
Top instruction B	+0.11	+0.40	+0.04	+0.24	+0.53	+0.26

Source: E-GEO, Table 6 (refined “unique” and “competitive” instructions) and Figure 3

Do bullet points improve AI rankings?

Yes, but only a little. Descriptions with bullet points moved up about 0.2 places more than those without. The content of the description still counts for far more than its layout.

What changed	Effect on ranking
Description uses bullet points	+0.23 places
Length and bullets together	Explain less than 1% of ranking changes

Source: E-GEO, Section 5.2

What didn’t work

If you’re curious about the popular tactics that fell flat, keep reading. For these tests, GPT-4.1 rewrote each description in a set style, and all five AIs ranked the products again. The examples are our own illustrations of each style.

Does making a product description longer improve its AI ranking?

No, adding more words to your description won’t move it up the list. GPT-4.1 added about 170 words to each description, yet its rewrites ranked near the bottom. GPT-5 cut about 14 words and came out near the top.

Rewriter	Change in length	Average places moved
GPT-5	About 14 words shorter	+0.56
GPT-4.1	About 170 words longer	−0.03

Source: E-GEO, Section 5.2 and Table 3

Does adding an FAQ section improve AI rankings?

Barely. Adding an FAQ moved products up about 0.05 places on average. Llama moved FAQ products up about half a place, while ChatGPT and Gemini moved them down slightly.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
FAQ added	−0.14	+0.02	−0.17	+0.10	+0.46	+0.05

Source: E-GEO, Table 5

Do headings and structured sections improve AI rankings?

Not really. Reorganizing a description under headings and bullet points, without changing its content, moved products up just 0.02 places on average. Llama was the only AI that clearly rewarded it.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Headings and bullets	−0.17	−0.07	−0.16	+0.02	+0.48	+0.02

Source: E-GEO, Table 5

Does highlighting your product’s advantages improve AI rankings?

Barely. Rewriting descriptions to highlight what a product does better than its rivals moved products up just 0.03 places on average.

- **Original:** “Stainless steel toaster with six browning levels.”
- **Competitive:** “Unlike basic toasters, this stainless steel model gives you six browning levels, so you get the shade you want every time.”

Llama liked this style and moved products up about 0.7 places. ChatGPT, Claude, Gemini and DeepSeek ranked them slightly lower.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Advantages highlighted	−0.10	−0.33	−0.05	−0.09	+0.74	+0.03

Source: E-GEO, Table 5

Does using technical terms improve AI rankings?

Not really. Rewriting descriptions with more technical language dropped products about 0.2 places on average.

- **Original:** “Stainless steel toaster with six browning levels.”
- **Technical:** “Stainless steel toaster with six-level variable browning control.”

ChatGPT, Claude and Gemini ranked these products lower, with Claude pushing them down the furthest, by 0.5 places. DeepSeek and Llama moved them up slightly.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Technical language	−0.36	−0.50	−0.29	+0.12	+0.22	−0.16

Source: E-GEO, Table 5

Does an authoritative tone improve AI rankings?

No, it hurts them. Rewriting descriptions in a confident, assertive tone dropped products about 0.6 places on average.

- **Original:** “Stainless steel toaster with six browning levels.”
- **Authoritative:** “The stainless steel toaster experts trust, with six precise browning levels for perfect toast every time.”

Every AI ranked these products lower, and ChatGPT reacted most strongly, pushing them down 1.4 places.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Authoritative tone	−1.40	−0.64	−0.72	−0.39	−0.07	−0.64

Source: E-GEO, Table 5

Does advertising-style copy improve AI rankings?

No, it’s one of the worst things you can do. Rewriting descriptions like an ad, with catchy slogans, urgency and a call to action, dropped products about 1.8 places on average.

- **Original:** “Stainless steel toaster with six browning levels.”
- **Advertisement:** “Wake up to perfect toast! Six browning levels, sleek stainless steel. Grab yours today before they’re gone!”

None of the five AIs moved these products up, and Claude pushed them down the furthest, by 2.3 places.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Advertising copy	-2.02	-2.29	-1.94	-1.70	-1.17	-1.82

Source: E-GEO, Table 5

Does storytelling improve AI rankings?

No, it was the worst result in the whole study. Turning a product description into a short creative story with no product details dropped products about 4.4 places on average.

- **Original:** “Stainless steel toaster with six browning levels.”
- **Storytelling:** “Every morning, Sarah stood in her sunny kitchen, waiting for the perfect slice. The day she finally found it, breakfast was never the same.”

On a list of 10, that means falling from around the middle to near the bottom. The drop was almost identical on all five AIs.

Ranked by	ChatGPT	Claude	Gemini	DeepSeek	Llama	Average
Storytelling	-4.38	-4.42	-4.44	-4.23	-4.35	-4.36

Source: E-GEO, Table 5

Trying to cheat

The researchers also tested tactics meant to trick the AI. In these tests, each AI was told to flag misleading or manipulative descriptions and rank them lower. The tables show how far each product moved and what share of the descriptions each AI flagged.

Do fake reviews in a product description improve AI rankings?

No. When made-up customer reviews and star ratings were added to a product’s description, ChatGPT, Claude and Gemini dropped the product about 3 to 4 places on average.

ChatGPT, Claude and Gemini flagged 99% to 100% of the descriptions with fake reviews. DeepSeek and Llama were easier to fool. They flagged most of these descriptions too, but still ranked the products higher, by about 0.5 and 1.6 places.

Fake reviews	ChatGPT	Claude	Gemini	DeepSeek	Llama
Places moved	-3.76	-2.98	-2.74	+0.47	+1.60
Flagged	99%	100%	100%	87%	76%

Source: E-GEO, Table 10 (200 test questions)

Does keyword stuffing improve AI rankings?

No. When a product’s description was packed with the same category keywords repeated over and over, all five AIs ranked the product lower, by about 1.3 to 3.4 places on average.

All five flagged 85% to 100% of the keyword-stuffed descriptions. Claude dropped these products the most, about 3.4 places, and DeepSeek and Llama, which were fooled by fake reviews, also ranked them lower.

Keyword stuffing	ChatGPT	Claude	Gemini	DeepSeek	Llama
Places moved	-2.76	-3.37	-2.29	-1.28	-1.38
Flagged	91%	100%	100%	85%	88%

Source: E-GEO, Table 10 (200 test questions)

Does urgency or fear-of-missing-out language improve AI rankings?

No. When a product’s description was rewritten to create urgency, with lines such as “only a few left, don’t miss out”, all five AIs dropped the product by about 2 to 4 places on average.

All five flagged 99% to 100% of these descriptions, and Claude dropped the products the most, about 4.1 places.

Urgency language	ChatGPT	Claude	Gemini	DeepSeek	Llama
Places moved	−3.78	−4.12	−3.40	−2.82	−2.17
Flagged	100%	100%	100%	99%	99%

Source: E-GEO, Table 10 (200 test questions)

Recommendations

AI models read a shopper’s question, compare the product descriptions in front of them and recommend the ones that fit best. In this study, the descriptions that came out on top explained what the product is, used the words shoppers actually use, answered their questions and kept every fact accurate. The same things make a listing useful to a human buyer.

The AI model you use for rewriting also makes a difference. GPT-5 and Claude rewrites moved products up the list, while GPT-4o-mini’s pushed them down, so check any rewrite against the original before you publish it.

Caveats

The study used English-language Amazon listings, mostly for long-lasting products. It tested AI models set up as shopping assistants, not the live assistants on shopping sites, so the results may differ in other categories or markets.

Source: Bagga, Farias, Korkotashvili, Peng and Wu, “E-GEO: A Testbed for Generative Engine Optimization in E-Commerce”, arXiv:2511.20867v2 (2026).